

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score, StratifiedKFold
from sklearn.preprocessing import StandardScaler, RobustScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_curve, auc, recall_score
from sklearn.pipeline import Pipeline
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as ImbPipeline
import warnings
warnings.filterwarnings('ignore')
```

```
# Load the data
df = pd.read_csv('/content/healthcare-dataset-stroke-data.csv')
```

```
# Data exploration and preprocessing
print("Original shape:", df.shape)
print("Missing values:\n", df.isnull().sum())
```

```
Original shape: (5110, 12)
Missing values:
  id          0
gender        0
age           0
hypertension  0
heart_disease 0
ever_married  0
work_type     0
Residence_type 0
avg_glucose_level 0
bmi          201
smoking_status 0
stroke        0
dtype: int64
```

```
# Feature engineering and improved preprocessing
# Handle missing values better - impute BMI instead of dropping
df['bmi'] = df['bmi'].fillna(df['bmi'].median())
```

```
# Drop the ID column
df = df.drop(columns=['id'])
```

```
# Check for 'Other' in gender and handle it
if 'Other' in df['gender'].values:
    # Either combine with majority class or create a new category
    df = df[df['gender'] != 'Other'] # Removing for simplicity as these are typically few
```

```
# Convert 'ever_married' to binary
df['ever_married'] = df['ever_married'].map({'Yes': 1, 'No': 0})
```

```
# Create age groups that might be more predictive
df['age_group'] = pd.cut(df['age'], bins=[0, 18, 35, 50, 65, 100],
                        labels=['Child', 'Young_Adult', 'Adult', 'Senior', 'Elderly'])
```

```
# Feature interaction: BMI × age might be important for stroke
df['bmi_age'] = df['bmi'] * df['age']
```

```
# Handling glucose levels - create categories
df['glucose_level'] = pd.qcut(df['avg_glucose_level'], q=4, labels=['Low', 'Medium', 'High', 'Very_High'])
```

```
# Better encoding
# One-hot encoding for categorical variables
cat_features = ['gender', 'work_type', 'Residence_type', 'smoking_status', 'age_group', 'glucose_level']
df_encoded = pd.get_dummies(df, columns=cat_features, drop_first=True)
```

```
# Separate features and target
X = df_encoded.drop('stroke', axis=1)
y = df_encoded['stroke']
```

```
# Print column names to verify encoding
print("\nFeatures after preprocessing:")
```

```
print(X.columns.tolist())
print(f"Total features: {X.shape[1]}")
```

Features after preprocessing:
['age', 'hypertension', 'heart_disease', 'ever_married', 'avg_glucose_level', 'bmi', 'bmi_age', 'gender_Male', 'work_type_Never_worked', 'work_type_Currently_working', 'residence_type_Urban', 'residence_type_Rural', 'average_yearly_salary']
Total features: 23

```
# Split data with stratification to maintain class distribution
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
print("Class distribution in training set:")
print(y_train.value_counts(normalize=True))
```

Class distribution in training set:
stroke
0 0.951309
1 0.048691
Name: proportion, dtype: float64

```
# Function to evaluate and print model performance
def evaluate_model(model, X_test, y_test, model_name="Model"):
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)[:, 1]

    # Calculate metrics
    acc = accuracy_score(y_test, y_pred)
    tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
    sensitivity = recall_score(y_test, y_pred, pos_label=1)
    specificity = tn / (tn + fp)

    print(f"\n{model_name} Results:")
    print(f"Accuracy    : {round(acc * 100, 2)}%")
    print(f"Sensitivity (Recall for 1): {round(sensitivity * 100, 2)}%")
    print(f"Specificity (Recall for 0): {round(specificity * 100, 2)}%")
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred))

    # Plot confusion matrix
    plt.figure(figsize=(6, 4))
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(f"Confusion Matrix - {model_name}")
    plt.show()

    # ROC curve
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)

    plt.figure(figsize=(6, 4))
    plt.plot(fpr, tpr, label=f'AUC = {roc_auc:.3f}')
    plt.plot([0, 1], [0, 1], 'k--')
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.title(f"ROC Curve - {model_name}")
    plt.legend()
    plt.show()

    return acc, sensitivity, specificity, roc_auc
```

```
# ----- APPROACH 1: Advanced Logistic Regression with proper scaling -----
# Create a pipeline with scaling and logistic regression
pipe_lr = Pipeline([
    ('scaler', StandardScaler()),
    ('classifier', LogisticRegression(max_iter=1000, class_weight='balanced', C=0.1))
])
```

```
# Train the model
pipe_lr.fit(X_train, y_train)
```

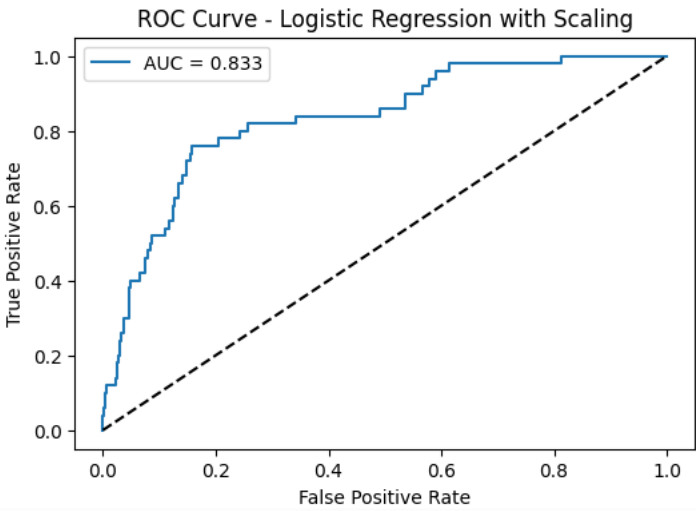

```
graph LR
    Pipeline[Pipeline] --> StandardScaler[StandardScaler]
    Pipeline --> LogisticRegression[LogisticRegression]
```

```
# Evaluate
lr_metrics = evaluate_model(pipe_lr, X_test, y_test, "Logistic Regression with Scaling")
```

Logistic Regression with Scaling Results:
Accuracy : 69.47%
Sensitivity (Recall for 1): 82.0%
Specificity (Recall for 0): 68.83%


Classification Report:

	precision	recall	f1-score	support
0	0.99	0.69	0.81	972
1	0.12	0.82	0.21	50
accuracy			0.69	1022
macro avg	0.55	0.75	0.51	1022
weighted avg	0.94	0.69	0.78	1022



```
# ----- APPROACH 2: Using SMOTE to handle class imbalance -----
# Create a pipeline with SMOTE
smote_pipe = ImbPipeline([
    ('scaler', StandardScaler()),
    ('smote', SMOTE(random_state=42)),
    ('classifier', LogisticRegression(max_iter=1000, C=0.1))
])
```

```
smote_pipe.fit(X_train, y_train)
```



```
# Evaluate
smote_metrics = evaluate_model(smote_pipe, X_test, y_test, "Logistic Regression with SMOTE")
```



Logistic Regression with SMOTE Results:

Accuracy : 69.47%

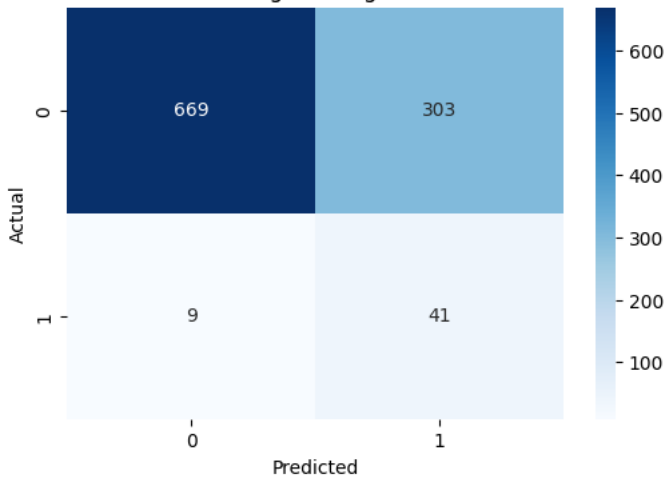
Sensitivity (Recall for 1): 82.0%

Specificity (Recall for 0): 68.83%


Classification Report:

	precision	recall	f1-score	support
0	0.99	0.69	0.81	972
1	0.12	0.82	0.21	50
accuracy			0.69	1022
macro avg	0.55	0.75	0.51	1022
weighted avg	0.94	0.69	0.78	1022

Confusion Matrix - Logistic Regression with SMOTE



ROC Curve - Logistic Regression with SMOTE



```
# ----- APPROACH 3: Random Forest with hyperparameter tuning -----
```

```
# Define the parameter grid
```

```
param_grid = {  
    'n_estimators': [100, 200],  
    'max_depth': [None, 10, 20],  
    'min_samples_split': [2, 5],  
    'class_weight': ['balanced', 'balanced_subsample']  
}
```

```
# Create RandomForest model
```

```
rf = RandomForestClassifier(random_state=42)
```

```
# Create grid search with cross-validation
```

```
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=cv, scoring='recall', n_jobs=-1)
```

```
# Train the model
```

```
grid_search.fit(X_train, y_train)
```



```
# Get the best model
best_rf = grid_search.best_estimator_
print(f"\nBest RF Parameters: {grid_search.best_params_}")
```



Best RF Parameters: {'class_weight': 'balanced', 'max_depth': 10, 'min_samples_split': 5, 'n_estimators': 100}

```
# Evaluate the best model
rf_metrics = evaluate_model(best_rf, X_test, y_test, "Optimized Random Forest")
```

Optimized Random Forest Results:
Accuracy : 91.1%
Sensitivity (Recall for 1): 34.0%
Specificity (Recall for 0): 94.03%

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.94	0.95	972
1	0.23	0.34	0.27	50
accuracy			0.91	1022
macro avg	0.60	0.64	0.61	1022
weighted avg	0.93	0.91	0.92	1022



```
# ----- APPROACH 4: Gradient Boosting model -----
gb = GradientBoostingClassifier(n_estimators=200, learning_rate=0.1,
                                max_depth=3, random_state=42)
gb.fit(X_train, y_train)
```



```
GradientBoostingClassifier
GradientBoostingClassifier(n_estimators=200, random_state=42)
```

```
# Evaluate
gb_metrics = evaluate_model(gb, X_test, y_test, "Gradient Boosting")
```

Gradient Boosting Results:
Accuracy : 94.62%
Sensitivity (Recall for 1): 4.0%
Specificity (Recall for 0): 99.28%


Classification Report:

	precision	recall	f1-score	support
0	0.95	0.99	0.97	972
1	0.22	0.04	0.07	50
accuracy			0.95	1022
macro avg	0.59	0.52	0.52	1022
weighted avg	0.92	0.95	0.93	1022



```
# ----- APPROACH 5: SVM with class weights and proper scaling -----
svm_pipe = Pipeline([
    ('scaler', RobustScaler()), # RobustScaler is good for datasets with outliers
    ('classifier', SVC(probability=True, class_weight='balanced', C=1.0, gamma='scale', kernel='rbf'))
])
```

```
# Train
svm_pipe.fit(X_train, y_train)
```





```
# Evaluate
svm_metrics = evaluate_model(svm_pipe, X_test, y_test, "Support Vector Machine")
```



Support Vector Machine Results:
Accuracy : 73.19%
Sensitivity (Recall for 1): 76.0%
Specificity (Recall for 0): 73.05%

Classification Report:


	precision	recall	f1-score	support
0	0.98	0.73	0.84	972
1	0.13	0.76	0.22	50
accuracy			0.73	1022
macro avg	0.56	0.75	0.53	1022
weighted avg	0.94	0.73	0.81	1022



```
# ----- APPROACH 6: Ensemble of models (Voting) -----
from sklearn.ensemble import VotingClassifier

# Create a voting classifier
voting_clf = VotingClassifier(
    estimators=[
        ('lr', pipe_lr),
        ('rf', best_rf),
        ('gb', gb),
        ('svm', svm_pipe)
    ],
    voting='soft' # Use probability estimates for prediction
)
```

```
# Train
voting_clf.fit(X_train, y_train)
```



```
# Evaluate
ensemble_metrics = evaluate_model(voting_clf, X_test, y_test, "Ensemble Model (Voting)")
```




Ensemble Model (Voting) Results:
Accuracy : 94.52%
Sensitivity (Recall for 1): 10.0%
Specificity (Recall for 0): 98.87%


Classification Report:

	precision	recall	f1-score	support
0	0.96	0.99	0.97	972
1	0.31	0.10	0.15	50
accuracy			0.95	1022
macro avg	0.63	0.54	0.56	1022
weighted avg	0.92	0.95	0.93	1022

Confusion Matrix - Ensemble Model (Voting)



ROC Curve - Ensemble Model (Voting)



```
# Compare all models
```

```
models = ["Logistic Regression", "LR with SMOTE", "Random Forest", "Gradient Boosting", "SVM", "Ensemble"]
metrics = [lr_metrics, smote_metrics, rf_metrics, gb_metrics, svm_metrics, ensemble_metrics]
```

```
# Create comparison dataframe
```

```
comparison = pd.DataFrame({
    'Model': models,
    'Accuracy (%)': [round(m[0]*100, 2) for m in metrics],
    'Sensitivity (%)': [round(m[1]*100, 2) for m in metrics],
    'Specificity (%)': [round(m[2]*100, 2) for m in metrics],
    'AUC': [round(m[3], 3) for m in metrics]
})
```

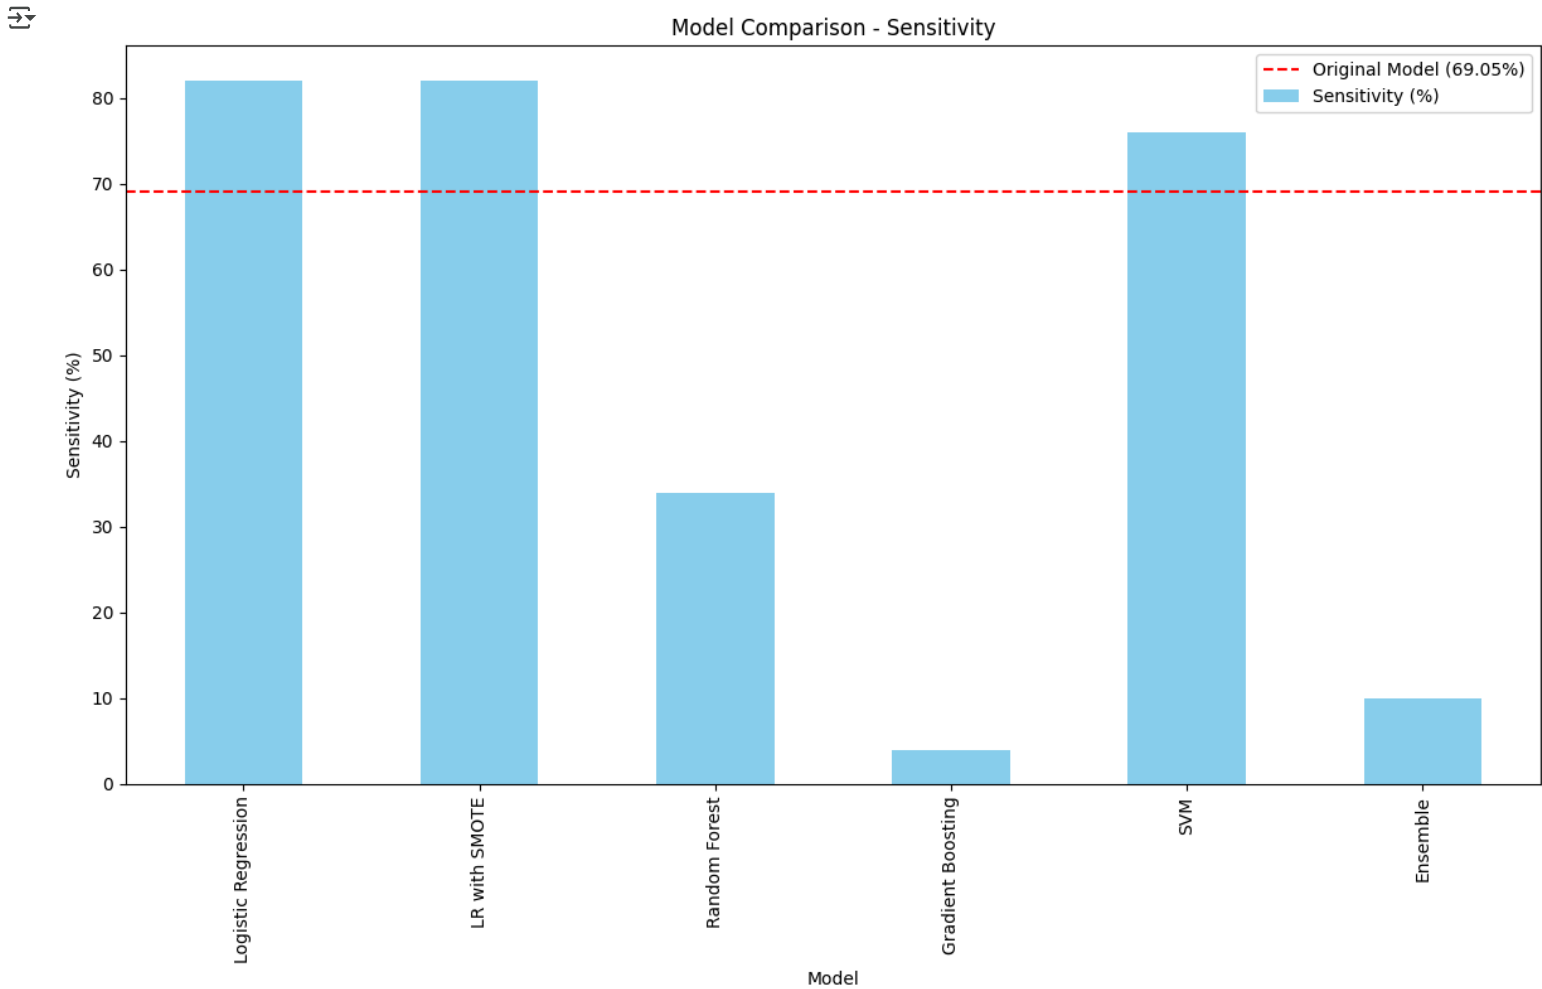


```
print("\n--- Model Comparison ---")
print(comparison)
```

↕

--- Model Comparison ---					
	Model	Accuracy (%)	Sensitivity (%)	Specificity (%)	AUC
0	Logistic Regression	69.47	82.0	68.83	0.833
1	LR with SMOTE	69.47	82.0	68.83	0.834
2	Random Forest	91.10	34.0	94.03	0.800
3	Gradient Boosting	94.62	4.0	99.28	0.787
4	SVM	73.19	76.0	73.05	0.792
5	Ensemble	94.52	10.0	98.87	0.826

```
# Visualize model performance
plt.figure(figsize=(12, 8))
comparison.set_index('Model')['Sensitivity (%)'].plot(kind='bar', color='skyblue')
plt.axhline(y=69.05, color='r', linestyle='--', label='Original Model (69.05%)')
plt.ylabel('Sensitivity (%)')
plt.title('Model Comparison - Sensitivity')
plt.legend()
plt.tight_layout()
plt.show()
```




```
# Feature importance from best model (if Random Forest or Gradient Boosting is best)
if 'feature_importances_' in dir(best_rf):
    importances = best_rf.feature_importances_
    indices = np.argsort(importances)[-10:]

    plt.figure(figsize=(12, 8))
    plt.title('Feature Importance')
    plt.bar(range(X.shape[1]), importances[indices], align='center')
    plt.xticks(range(X.shape[1]), X.columns[indices], rotation=90)
    plt.tight_layout()
    plt.show()

    print("\nTop 10 Most Important Features:")
    for i in range(min(10, X.shape[1])):
        print(f"{i+1}. {X.columns[indices[i]]}: {importances[indices[i]]:.4f}")
```



Feature Importance



Top 10 Most Important Features:

- 1. bmi_age: 0.2323
- 2. age: 0.2183
- 3. bmi: 0.1145
- 4. avg_glucose_level: 0.1038
- 5. age_group_Elderly: 0.0865
- 6. age_group_Young_Adult: 0.0293
- 7. hypertension: 0.0242
- 8. ever_married: 0.0211